

Introduction to SQL

[View All Tutorials](#)[Login to Run Tutorial](#)

Tutorial	Introduction to SQL
Description	This tutorial provides an introduction to the Structured Query Language (SQL), learn how to create tables with primary keys, columns, constraints, indexes, and foreign keys.
Tags	create table, create, select, insert, update, delete, drop, drop table, recycle bin, purge
Area	SQL General
Contributor	Mike Hichwa (Oracle)
Created	Monday October 12, 2015
Modules	13

Module

1

Creating Tables

Tables are the basic unit of data storage in an Oracle Database. Data is stored in rows and columns. You define a table with a table name, such as employees, and a set of columns. You give each column a column name, such as employee_id, last_name, and job_id; a datatype, such as VARCHAR2, DATE, or NUMBER; and a width. The width can be predetermined by the datatype, as in DATE. If columns are of the NUMBER datatype, define precision and scale instead of width. A row is a collection of column information corresponding to a single record.

You can specify rules for each column of a table. These rules are called integrity constraints. One example is a NOT NULL integrity constraint. This constraint forces the column to contain a value in every row.

For example:

```
create table DEPARTMENTS (  
  deptno      number,  
  name        varchar2(50) not null,  
  location    varchar2(50),  
  constraint pk_departments primary key (deptno)  
);
```

Tables can declarative specify relationships between tables, typically referred to as referential integrity. To see how this works we can create a "child" table of the DEPARTMENTS table by including a foreign key in the EMPLOYEES table that references the DEPARTMENTS table. For

Inserting Data

Now that we have tables created, and we have triggers to automatically populate our primary keys, we can add data to our tables. Because we have a parent child relationship, with the DEPARTMENTS table as the parent table, and the EMPLOYEES table as the child we will first INSERT a row into the DEPARTMENTS table.

```
insert into departments (name, location) values
('Finance', 'New York');
```

```
insert into departments (name, location) values
('Development', 'San Jose');
```

Lets verify that the insert was successful by running a SQL SELECT statement to query all columns and all rows of our table.

```
select * from departments;
```

You can see that an ID will have been automatically generated. You can now insert into the EMPLOYEES table a new row but you will need to put the generated DEPTID value into your SQL INSERT statement. The examples below show how we can do this using a SQL query, but you could simply enter the department number directly.

```
insert into EMPLOYEES
(name, job, salary, deptno)
values
('Sam Smith', 'Programmer',
5000,
(select deptno
from departments
where name = 'Development'));
```

```
insert into EMPLOYEES
(name, job, salary, deptno)
values
('Mara Martin', 'Analyst',
6000,
(select deptno
from departments
where name = 'Finance'));
```

```
insert into EMPLOYEES
(name, job, salary, deptno)
values
('Yun Yates', 'Analyst',
5500,
(select deptno
from departments
where name = 'Development'));
```

Typically developers index columns for three major reasons:

1. To enforce unique values within a column
2. To improve data access performance
3. To prevent lock escalation when updating rows of tables that use declarative referential integrity

When a table is created and a PRIMARY KEY is specified an index is automatically created to enforce the primary key constraint. If you specify UNIQUE for a column when creating a column a unique index is also created. To see the indexes that already exist for a given table you can run the following dictionary query.

```
select table_name "Table",
       index_name  "Index",
       column_name "Column",
       column_position "Position"
from   user_ind_columns
where  table_name = 'EMPLOYEES' or
       table_name = 'DEPARTMENTS'
order by table_name, column_name, column_position
```

It is typically good form to index foreign keys, foreign keys are columns in a table that reference another table. In our EMPLOYEES and DEPARTMENTS table example the DEPTNO column in the EMPLOYEE table references the primary key of the DEPARTMENTS table.

```
create index employee_dept_no_fk_idx
on employees (deptno)
```

We may also determine that the EMPLOYEE table will be frequently searched by the NAME column. To improve the performance searches and to ensure uniqueness we can create a unique index on the EMPLOYEE table NAME column.

```
create unique index employee_ename_idx
on employees (name)
```

Oracle provides many other indexing technologies including function based indexes which can index expressions, such as an upper function, text indexes which can index free form text, bitmapped indexes useful in data warehousing. You can also create indexed organized tables, you can use partition indexes and more. Sometimes it is best to have fewer indexes and take advantage of in memory capabilities. All of these topics are beyond the scope of this basic introduction.

Querying Data

To select data from a single table it is reasonably easy, simply use the SELECT ... FROM ... WHERE ... ORDER BY ... syntax.

```
select * from employees;
```

To query data from two related tables you can join the data

```
select e.name employee,
       d.name department,
       e.job,
       d.location
from departments d, employees e
where d.deptno = e.deptno(+)
order by e.name;
```

As an alternative to a join you can use an inline select to query data.

```
select e.name employee,
       (select name
        from departments d
        where d.deptno = e.deptno) department,
       e.job
from employees e
order by e.name;
```

Module

6

Adding Columns

You can add additional columns after you have created your table using the ALTER TABLE ... ADD ... syntax. For example:

```
alter table EMPLOYEES
add country_code varchar2(2);
```

Module

7

Querying the Oracle Data Dictionary

Table meta data is accessible from the Oracle data dictionary. The following queries show how you can query the data dictionary tables.

```
select table_name, tablespace_name, status
from user_tables
where table_name = 'EMPLOYEES';
```

```
select column_id, column_name, data_type
from user_tab_columns
where table_name = 'EMPLOYEES'
order by column_id;
```

Module

8

Updating Data

You can use SQL to update values in your table, to do this we will use the update clause

```
update employees
set country_code = 'US';
```

The query above will update all rows of the employee table and set the value of country code

to US. You can also selectively update just a specific row.

```
update employees
set commission = 2000
where name = 'Sam Smith';
```

Lets run a Query to see what our data looks like

```
select name, country_code, salary, commission
from employees
order by name;
```

Module

9

Aggregate Queries

You can sum data in tables using aggregate functions. We will use column aliases to rename columns for readability, we will also use the null value function (NVL) to allow us to properly sum columns with null values.

```
select
    count(*) employee_count,
    sum(salary) total_salary,
    sum(commission) total_commission,
    min(salary + nvl(commission,0)) min_compensation,
    max(salary + nvl(commission,0)) max_compensation
from employees;
```

Module

10

Compressing Data

As your database grows in size to gigabytes or terabytes and beyond, consider using table compression. Table compression saves disk space and reduces memory use in the buffer cache. Table compression can also speed up query execution during reads. There is, however, a cost in CPU overhead for data loading and DML. Table compression is completely transparent to applications. It is especially useful in online analytical processing (OLAP) systems, where there are lengthy read-only operations, but can also be used in online transaction processing (OLTP) systems.

You specify table compression with the COMPRESS clause of the CREATE TABLE statement. You can enable compression for an existing table by using this clause in an ALTER TABLE statement. In this case, the only data that is compressed is the data inserted or updated after compression is enabled. Similarly, you can disable table compression for an existing compressed table with the ALTER TABLE...NOCOMPRESS statement. In this case, all data the was already compressed remains compressed, and new data is inserted uncompressed.

To enable compression for future data use the following syntax.

```
alter table EMPLOYEES compress for oltp;
alter table DEPARTMENTS compress for oltp;
```

Deleting Data

You can delete one or more rows from a table using the DELETE syntax. For example to delete a specific row:

```
delete from employees
where name = 'Sam Smith';
```

Dropping Tables

You can drop tables using the SQL DROP command. Dropping a table will remove all of the rows and drop sub-objects including indexes and triggers. The following DROP statements will drop the departments and employees tables. The optional cascade constraints clause will drop remove constraints thus allowing you to drop database tables in any order.

```
drop table departments cascade constraints;
drop table employees cascade constraints;
```

Un-dropping Tables

If the RECYCLEBIN initialization parameter is set to ON (the default in 10g), then dropping this table will place it in the recycle bin. To see if you can undrop a table run the following data dictionary query:

```
select object_name,
       original_name,
       type,
       can_undrop,
       can_purge
from recyclebin;
```

To undrop tables we use the flashback command, for example:

```
flashback table DEPARTMENTS to before drop;
flashback table EMPLOYEES to before drop;
select count(*) departments
from departments;
select count(*) employees
from employees;
```

Additional Information

Database on OTN

[SQL and PL/SQL Discussion forums](#)

[Oracle Database](#)

[Download Oracle Database](#)

[Latest Database Tutorials](#)

ORACLE®

Integrated Cloud
Applications & Platform Services

© 2018 Oracle Corporation

[Privacy](#) | [Terms of Use](#)